

# Data Modelling and Databases

Fabian Bösigler

## Contents

Relational Model . . . . .	1
Database Schema . . . . .	1
Relation Schema . . . . .	1
Fields/Attributes . . . . .	2
Relational Algebra . . . . .	2
Relational Calculus . . . . .	3
SQL . . . . .	3
Create Table . . . . .	3
Integrity Constraints . . . . .	4
Recursion . . . . .	4
Functional Dependencies . . . . .	4
Normal Form . . . . .	4
Entity Relationship Model . . . . .	4
Entity Relationship to Relational Model . . . . .	4
Database System Overview . . . . .	4
Transactions and ACID . . . . .	4
Isolation and Locking . . . . .	4
Recovery . . . . .	5

## Relational Model

### Database Schema

A set of relation schemas.

### Relation Schema

A name and a set of fields/attributes.

For a relation  $R(f_1 : D_1, \dots, f_n : D_n)$ , an Instance  $I_R$  is a set of tuples  $I_R \subseteq D_1 \times \dots \times D_n$ .

## Fields/Attributes

A name and a domain (e.g. Integer, String).

## Relational Algebra

Relational algebra is an imperative way of obtaining the data. It is step-by-step instructions on how to obtain the result.

## Basic Operators

- Union:  $\cup$
- Difference:  $-$
- Selection:  $\sigma$
- Projection:  $\Pi$
- Cartesian Product:  $\times$
- Renaming:  $\rho$

## Joins

**Natural Join**  $R_1(A, B) \bowtie R_2(B, C) = \Pi_{A, B, C}(\sigma_{R_1.B=R_2.B}(R_1 \times R_2))$

If  $R_1$  and  $R_2$  have no shared attributes,  $R_1 \bowtie R_2 = R_1 \times R_2$ .

If  $R_1$  and  $R_2$  share all attributes,  $R_1 \bowtie R_2 = R_1 \cap R_2$ .

**Theta Join**  $R_1 \bowtie_{\theta} R_2 = \sigma_{\theta}(R_1 \times R_2)$

**Equi-Join**  $R_1 \bowtie_{A=B} R_2 = \sigma_{A=B}(R_1 \times R_2)$

**Semi-Join**  $R_1(A_1, \dots, A_n) \bowtie_C R_2(B_1, \dots, B_m) = \Pi_{A_1, \dots, A_n}(R_1 \bowtie R_2)$

**Relational Division**  $R \div S = T$ , where  $T$  is the largest relation such that  $S \times T \subseteq R$ .

$$R \div S = \Pi_{R-S}R - \Pi_{R-S}((\Pi_{R-S}R) \times S - R)$$

## Relational Calculus

Relational Calculus queries data in a declarative way. It tells the system what we want instead of how to get it.

Example: Get all students who take the database course.

$\Pi_{\text{PersNr}} \sigma_{\text{Title}=\text{"Database"}}((\text{Student} \bowtie \text{Attends}) \bowtie \text{Lecture})$

Is equivalent to:

$\{p : \exists n, s. \text{Student}(p, n, s) \wedge \exists l, t, c, r. (\text{Attends}(p, l) \wedge \text{Lecture}(l, t, c, r) \wedge t = \text{"Database"})\}$

**Safety** We say a relational calculus query  $Q$  is safe, if  $Q(I)$  is finite for all instances  $I$ . An unsafe example would be  $\{x \mid \neg R(x)\}$ . The safety problem is undecidable.

## SQL

### Create Table

```
CREATE TABLE professor (  
    pers_nr INTEGER,  
    name    VARCHAR(32),  
    level   CHARACTER(2) DEFAULT "AP",  
    PRIMARY KEY (pers_nr)  
)
```

TODO

**Integrity Constraints**

**Recursion**

**Functional Dependencies**

**Normal Form**

**Entity Relationship Model**

**Entity Relationship to Relational Model**

**Database System Overview**

**Transactions and ACID**

**Isolation and Locking**

Every transaction starts with **BEGIN** and ends with either **COMMIT** or **ABORT**. If **COMMIT**, all changes are saved, if **ABORT**, all changes are undone.

**Serializability Classes**

**Serializable** Schedules that lead to the same answer as some serial schedules.

		10, 10
Read(A, t)		
t := t + 100		
Write(A, t)		110, 10
	Read(A, s)	
	s := s * 2	
	Write(A, s)	220, 10
Read(B, t)		
t := t + 100		
Write(B, t)		220, 110
	Read(B, s)	
	s := s * 2	
	Write(B, s)	220, 220

**Conflict Serializable** We can translate a schedule into a serial shedule with a sequence of nonconflicting swaps of adjacent actions.

Types of Conflicts:

- Read-Write Conflict
- Write-Read Conflict
- Write-Write Conflict

## Locking

### Two-Phase Locking Growing Phase:

- Each transaction requests the locks that it needs
- Locks cannot be released in this phase

### Shrinking Phase:

- The transaction is only allowed to release locks that it previously required
- The transaction cannot acquire new locks

### Cascading Abort

### Serial Transactions are serial.

	10, 10
Read(A, t)	
t := t + 100	
Write(A, t)	110, 10
Read(B, t)	
t := t + 100	
Write(B, t)	110, 110
Read(A, s)	
s := s * 2	
Write(A, s)	220, 110
Read(B, s)	
s := s * 2	
Write(B, s)	220, 220

## Recovery

### Recoverability Classes

**Recoverable** If  $T_i$  reads from  $T_j$  and commits, then  $c_j < c_i$ .

**Avoids Cascading Aborts** If  $T_i$  reads  $X$  from  $T_j$ , then  $c_j < r_i[X]$ .

**Strict** If  $T_i$  reads from or overwrites  $X$  written by  $T_j$ , then  $(c_j < r_i[X] \wedge c_j < w_i[X]) \vee (a_j < r_i[X] \wedge a_j < w_i[X])$ .